WEEK 2 · LECTURE L3 OF 3 · UNIT I

Round 1: The Diagnosis

Where do we start? The four decisions of Round 1, the anchor test, and the method of a sound diagnosis.

Part 1 The Round 1 decisions and the Evidence Rationale

Part 2 The anchor stress test: the corporate-first plan

Part 3 Method: from symptom to binding constraint

Read Your pod's Round 1 submission

R K Singh · Senior Professor, Department of Commerce · Delhi School of Economics, University of Delhi · DSEH001 · M.Com. PGCF 2025

HOUR THREE

R1 The Diagnosis — Round Briefing

0:00 The decision set · 0:25 The anchor stress test · 0:45 R1 opens

0:00 – 0:06 · READ THIS TWICE

Three currencies. One grant. Nothing refills.

₹6 crore

Budget

Capital, licenses and change management for all six rounds. Overspending is blocked; starving adoption is penalized.

100

Change Capital

Opens at 100, ceiling 150. Disruptive options spend it; honest communication earns it back. Below 30, everything gets harder.

100

Analyst Days

A single grant for the entire mandate. It never refills. No pod can buy analytical depth everywhere.

Scored on five equal axes to 20

Agility & Readiness · Financial Discipline · Employee Centricity · Analytics Capability · Equity & Compliance

Submission

Only the round's Pod Lead may submit. The analysis and Evidence Rationale autosave as anyone types — shared work, single submission. Rounds lock irreversibly.

0:06 – 0:18

R1 asks you for four decisions

01	**Which maturity model do you adopt?**	The model determines what you can see. Justify the choice against the diagnostic task — not as a preference.
02	**Which drivers dominate?**	Rank them and evidence the ranking. The trap: ranking technological first because technology is the visible subject.
03	**What is the strategic-alignment gap?**	One paragraph. Test it against outcome, ownership, sacrifice and reach. A gap statement naming no workforce segment is too abstract to use.
04	**Which single domain do you transform first?**	The hardest decision. Everything is broken and the budget is finite.

There is no single correct domain. The scorecard rewards a defended choice with honest trade-offs — and penalizes an undefended one, whichever domain it names.

Three defensible ways to pick the first domain

Foundational dependency

Does everything else require this first?

Records and data governance sit here.

Unglamorous — and frequently correct.

Pain and reach

Where is the largest measurable cost against the largest segment?

Frontline attrition and time-to-hire sit here.

Credibility and momentum

What earns Change Capital and buys the mandate for harder later moves?

A visible frontline win sits here.

Candidate domains visible in the pack: core records and the system of record · frontline hiring funnel · onboarding · learning · engagement measurement · the analytics function · data consent and governance ahead of DPDP

Name your criterion in the rationale. A choice without a stated criterion is a preference.

What your Evidence Rationale must contain

Claim	Is it specific enough to be wrong?
Evidence	Does it point to the pack or the workbook — with a location, not a memory?
Framework	Is a named model doing analytical work, or is it decoration?
Trade-off	What did this cost — budget, Change Capital, Analyst Days, or foregone domains?
Segment reach	Which of the four anchors does this serve, and which does it not?

A rationale with no trade-off in it is not a rationale.

It is a wish.

The Financial Discipline axis exists to find wishes.

0:28 – 0:40 · IN YOUR POD, EIGHT MINUTES

The Corporate-First Plan

A consultant proposes: Year 1, implement a cloud HRIS for the 2,400 head-office and regional staff. Clean the data, prove the model, build credibility with the Board. Extend to stores in Year 2, and to the gig network “when the regulatory position clarifies.”

It is cheap. It is low-risk. It will succeed. It is exactly what a competent consultant would recommend.

Take each of your four anchors in turn. Then decide: endorse, amend, or reject — and be ready to defend it.

0:40 – 0:45 · DEBRIEF, ONE ANCHOR PER POD

What the plan does to each anchor

GL-1001

Lakshmi Naidu

Nothing, for a year. The Year-2 extension inherits a system designed around laptop, English and company email.

Sequencing is architecture. What you defer, you constrain.

GL-1002

Arjun Joshi

Explicitly excluded pending regulatory clarity — which may arrive as an obligation, not a signal.

“When it clarifies” is a decision, not a deferral.

GL-1003

Sunil Mehta

In scope, with no role, no reskilling path, no reason to cooperate — and he holds the process knowledge.

Change Capital is spent or earned here.

GL-1004

Shreya Iyer

Best served. Genuinely — which is why the plan is seductive.

Serving the ready is not wrong. Serving only the ready is.

THE VERDICT TO LAND

The corporate-first plan is not obviously bad.

It is cheap, low-risk, will succeed on its own terms, and versions of it are implemented across Indian industry every year.

What it is, is a plan whose success metric has been quietly redefined to exclude three-quarters of the workforce.

The honest position: adopt corporate-first sequencing only if you can state, in the same breath, what you are doing in Year 1 to keep the Year-2 frontline extension architecturally possible and politically funded.

0:45 – 1:00

R1 opens. What happens next.

Agree now, on the platform

- Who drafts which section — the five roles map onto the four decisions plus the anchor check
- Your maturity model, provisionally, in writing
- Two workbook analyses and their cost in Analyst Days
- Your meeting point before Week 3, L1

The Week 3 cold-calls are on your pod's actual R1 decision — and any member may be called, not the Pod Lead.

R1 opens	Now — end of Week 2, L3
R1 closes	End of Week 3 · locks irreversibly
Pod charter	End of Week 2
R1 cold-calls	Week 3, L1 · first eight minutes
Cross-pod debrief	Week 3, L3

Week 3 pre-reading: Caldwell (2022) data-literacy and BI chapters · Manuti & de Palma critical chapter · DPDP orientation note

You now hold a diagnosis you will live with for eleven weeks.

In Week 4 you will architect a stack on top of it.

In Week 6 the adoption data will tell you whether the architecture was right.

In Week 8 the fairness audit will tell you what your diagnosis failed to look at.

Everything downstream inherits this week.

Diagnose honestly.

SESSION MAP

The hour ahead

60 min

0:00

From pattern to cause

What a diagnosis is, and the differential method

0:16

What counts as evidence

Four kinds, triangulation, and the contradiction

0:30

Which constraint binds

Binding and slack, four tests, sequencing

0:48

Error and craft

How diagnoses go wrong, and how to write one

0:57

Limits and close

What a correct diagnosis still cannot do

Leave with this: the ability to state a cause, name the evidence that discriminated in its favor, and say what would prove you wrong.

OPENING SOCRATIC DIALOGUE

Three findings. All true. Which one do you act on?

FINDING A

Workforce planning is 8% automated. People analytics is 10%. The two things the Board is asking for are the two the function cannot do.

FINDING B

Employees whose only device is a feature phone cannot reach the self-service app or the learning platform at all.

FINDING C

Five of the sixteen systems hold employee data outside India, and nobody has documented why.

The question to sit with

Your Round 1 rationale picked one, and your Round 2 architecture spent ₹120L acting on it. **What decided the order?** Not severity — all three are serious. Not cost — you did not know it yet. Something else was doing the work, and today we name it.

THE ANATOMY OF A DIAGNOSIS

Four things that are routinely confused

Symptom

What the organization complains of. “Our HR reporting is slow.” Reported, not observed — and the complaint may not be about the thing that is wrong.

Sign

What you can observe independently of anyone's account. Three systems return three headcounts. A sign does not require the organization to agree with you.

Pattern

Signs that occur together in a recognizable configuration. This is exactly what a framework supplies — “this is Level 2” names a pattern.

Cause

The mechanism that generates the pattern. Not “the systems are fragmented” but *why nobody stopped them fragmenting*.

Where last week ended, and this week begins

A framework takes you as far as **pattern**. It stops there by design. **A diagnosis is a claim about cause** — and a claim needs evidence that could have gone the other way.

THE METHOD

Eliminate rather than confirm

The instinct is to find a plausible cause and look for support. The discipline is the reverse: **list the causes that could have produced these signs, then look for the evidence that tells them apart.**

Step	What you do	Why it matters
1	**Fix the signs** you are explaining	A moving target can be explained by anything. Write the observations down before you theorize.
2	**List the candidate causes** — at least three	Two candidates is usually one candidate and a straw man. The third is where the honesty lives.
3	**Find the discriminating evidence**	The test worth running is the one whose result differs depending on which cause is true.
4	**State what survived, and what would kill it**	A diagnosis you cannot falsify is an opinion wearing a framework's clothes.

The discriminating test, in one line: if the evidence you are about to gather would look the same whichever cause is true, it is not worth gathering. Most analysis fails here, not at the conclusion.

WHAT COUNTS AS EVIDENCE

Four kinds, each strong for something different

Kind	Strong evidence about	Weak or silent about
System trace logins, transactions, tickets	What people actually did. Nobody performed it for your benefit.	Why they did it, and everyone the system never touched.
Administrative record HR fields, licenses, contracts	What the organization has committed to and paid for.	Reality. It records what was entered, not what is true.
Self-report surveys, interviews	Perception, intention, and what people are willing to say.	Behavior. People misremember, and they answer the question they think you asked.
Assertion “the reason is…”	Generating candidate causes. Insiders know where the bodies are.	Settling between them. An assertion is a hypothesis, never a finding.

This is not a ranking from best to worst. **It is a map of what each kind can and cannot carry.** Using self-report to establish behavior, or a system trace to establish motive, is the commonest evidential error in applied HR work.

TRIANGULATION

Three sources agreeing proves less than you think

Triangulation is the practice of establishing a finding from more than one kind of evidence. Its power comes entirely from **independence** — and independence is easy to lose without noticing.

False triangulation

The HR dashboard, the monthly report and the Board pack all show the same headcount. Three sources, one origin — they are three views of a single extract. If it is wrong, all three are wrong together.

Real triangulation

The system trace says few people complete a task. The helpdesk log says many people ask how. A supervisor says they do it on people's behalf. Three origins, three mechanisms, one conclusion.

The test to apply before you rely on it

Could these sources fail together? If a single upstream error, a single definition, or a single person could make all of them wrong at once, you have one source with three names.

THE MOST USEFUL THING YOU WILL FIND

A contradiction is data, not a nuisance

When two credible sources disagree, something structural is producing the disagreement. The disagreement is usually more informative than either source alone.

Definitional

Two true numbers counting different populations. Points at an absent shared definition.

Temporal

Two true numbers from different moments. Points at reconciliation that happens manually, or not at all.

Behavioral

What people say they do against what the trace shows. Points at a workaround somebody built.

Evaluative

Universal use alongside universal dislike. Points at something with no alternative.

WORKED ILLUSTRATION

Adoption 1.00, satisfaction 1.5

In the GreenLeaf inventory, spreadsheet trackers are used by everyone and rated lowest of any system. The contradiction is the finding: **spreadsheets are not a preference, they are a symptom.** Remove them without replacing the function they serve and they reappear under new file names.

THIS MOVE WAS MADE IN THIS ROOM

One pod read the spreadsheet as a symptom

SYS-16 — SPREADSHEET TRACKERS

1.00
adoption — the highest in the estate

1.5
satisfaction — the lowest in the estate

> *“SYS-16 is the one with highest adoption rate of 1 but lowest satisfaction score. It can be because employees are forced to use manual excel workarounds to fill the functional gaps between disconnected softwares.”*
>
> **A pod’s Round 2 rationale — the only pod to resolve it**

Ignore it
A spreadsheet tracker looks like furniture in a sixteen-system inventory.

Reconcile it away
Treat the satisfaction score as noise, or as people grumbling about a tool that works.

Explain it
Ask what state of the world produces *both* numbers at once. Only this one is a diagnosis.

The causal arrow reverses. The spreadsheet is not a cause of fragmentation; it is evidence of it. Ban it and the bridge goes while the gap stays — so it returns under new file names on a personal drive, where you can no longer see it.

Quoted verbatim from a submitted Round 2 rationale · 9

THE TURN OF THE HOUR

A diagnosis is not a list of what is wrong.
It is a claim about what, if fixed, changes the rest.

Everything so far has been about establishing what is true. Everything from here is about a different question, and a harder one: of the true things, which one is holding the others in place?

BINDING AND SLACK

You already know this idea from optimization

A binding constraint

Holds the solution back. Relax it by one unit and the objective improves. In a linear program it is the constraint with a non-zero shadow price.

A slack constraint

Real, correctly stated, and not currently limiting anything. Relax it and nothing moves. Effort spent here is genuinely wasted, however serious the constraint sounds.

Goldratt made the same point about factories: the throughput of a system is set by its bottleneck. **An hour saved at a non-bottleneck is not an hour saved — it is an hour of extra inventory.** Improvement away from the constraint converts into cost, not output.

The translation into HR

Fixing a problem that is not binding produces activity, cost, and a report saying the problem was addressed — and no change in any outcome anyone cares about. **It is indistinguishable from progress until the following year.**

FOUR TESTS

Is this constraint actually binding?

TEST 1

Relaxation

If this were removed tomorrow at no cost, name the specific thing downstream that would change. If you cannot name one, it is slack.

TEST 2

Precedence

Does anything else wait on this? Last week's six frameworks each supply a precedence rule — this is where you cash them in.

TEST 3

Recurrence

Has it already defeated one attempt? A constraint that has absorbed a previous remedy is usually real, and usually deeper than the remedy assumed.

TEST 4

Ownership

Is there someone who could actually relax it? A constraint nobody can move is a *condition* — plan around it, do not target it.

Tests 1 and 2 establish that it matters. Tests 3 and 4 establish that it is **tractable**. A diagnosis that passes the first two and fails the fourth is correct and useless — and correct-and-useless is the most common way good analysis dies.

CASHING IN LAST WEEK'S FRAMEWORKS

Six precedence rules make a partial order

Framework	The rule it contributes	What it forbids
Kotter	Readiness before rollout	Launching into an organization that does not believe there is a problem
Kavanagh	Foundation before capability	Buying a stage above the one you occupy
Bersin	Data before question	Asking a question the records cannot honestly answer
Westerman	Direction alongside investment	Spending faster than you can govern
Digital divide	Access, then skills, then use	Training people the system cannot reach
DPDP & SS2020	Not a step in the sequence at all	Treating a legal obligation as one option among others

Obligations are boundaries, not steps. They do not compete for a place in the order — they define the region inside which any order is permissible. A plan that sequences compliance after convenience has misunderstood what kind of thing it is.

WHEN THE ORDER IS STILL OPEN

Two constraints both bind. Now what?

If one precedes the other

Sequence them, and say why. The precedence rules settle it, and the reasoning belongs in the plan rather than in your head.

If they are independent

Run them in parallel — but only if they draw on different owners, different budgets and different attention. Two projects sharing one overstretched manager are not parallel.

If neither, use reversibility

Do the cheap-to-undo one first and learn from it. Defer the one that is expensive or humiliating to reverse until the uncertainty has come down.

The tie-breaker nobody teaches

Prefer the move that **makes the next diagnosis easier**. Fixing the record first does not only fix the record — it makes every subsequent question answerable. Some repairs pay twice.

HOW DIAGNOSES GO WRONG

Four failures, all of them comfortable

Premature closure

Stopping at the first cause that fits. The evidence gathered afterward is selected to confirm, and it always succeeds.

Solution-first

The remedy is already chosen — a product, a vendor, a restructure — and the diagnosis is reverse-engineered to require it.

The average

A mean reported over a population that has two distinct halves. The number is arithmetically correct and describes nobody.

Infinite regress

Asking why five times always arrives at leadership or culture. True, unfalsifiable, and unactionable.

Where to stop asking why

At the **last cause someone in the room can act on**. Going one step further feels more rigorous and delivers less: “the culture is not data-driven” is a sentence no budget can be written against.

THE WHOLE METHOD, ON ONE CASE

A worked differential

THE SIGNS

Sixteen systems in the inventory, several standing alone. Two systems of record running in parallel. Payroll mature; workforce planning at 8% automation. A live flight-risk model scoring every employee. Several processes with no named owner.

THREE CANDIDATE CAUSES

A Under-investment in technology. **B** No single record and no process ownership. **C** A skills deficit in the HR function.

WHAT DISCRIMINATES

If **A** were binding we would expect few systems. There are sixteen, and a fresh one went live last year. If **C** were binding we would expect even simple systems to fail. Payroll runs at the highest maturity in the estate and the highest satisfaction.

WHAT SURVIVES

B. The firm can buy and can operate. What is absent is one agreed record and a named owner with authority to impose it. That is why buying more has not helped.

What would falsify it: find a process that has one owner, one clean record, and still fails. If several exist, cause B is not carrying the weight this diagnosis puts on it.

THE CRAFT

Six moves of a defensible diagnosis

1 The observation — stated with its source, so a reader can go and look.

2 The pattern — and the framework that names it, cited as the course names it.

3 The candidates — the causes you considered, including the ones you rejected.

4 The discrimination — the evidence that separated them, and why it separates.

5 The claim — the surviving cause, with the consequence that follows for sequencing.

6 The falsifier — what you would accept as proof that you are wrong.

The one-sentence form

"Because *[observation, sourced]*, which is a case of *[pattern]*, and because *[discriminating evidence]* rules out *[rejected cause]*, the binding constraint is *[cause]* — so *[what must happen first]*."

If it will not fit this shape, one of the six moves is missing. Usually it is the third.

THE HONEST BOUNDARY

What a correct diagnosis still cannot do

Tell you it is affordable

Identifying the binding constraint says nothing about whether relaxing it is worth what it costs. That is a separate argument.

Predict resistance

The correct first move is often the one that takes authority away from somebody. Diagnosis is silent about who will fight it.

Stay true

It is a snapshot. Constraints move — and relaxing one promotes whatever was second into being the binding one.

See what it cannot measure

The analysis covers the system that produces data. Anyone the systems never touch is invisible to it by construction.

And the plainest limit of all: **a correct diagnosis acted on badly still fails.** Knowing what to do first is necessary and nowhere near sufficient — which is why change management is a separate discipline and not an afterthought.

Going deeper: the theory underneath diagnosis

The next slides set out Goldratt's focusing steps, abductive reasoning, confirmation bias and the premortem, and double-loop learning: what each claims, how it works, where it breaks, and what it means at GreenLeaf.

THEORY IN DEPTH · GOLDRATT (1984)

Every system has one binding constraint

THE CORE IDEA

The Theory of Constraints: output is limited by one bottleneck. Improving anything else changes little.

HOW IT WORKS

Five steps: identify the constraint, exploit it, subordinate everything else, elevate it, then repeat.

LIMITS AND CRITIQUE

Real organizations may have several interacting constraints, and some are policies, not machines.

AT GREENLEAF If the binding constraint is access, then training first is improving a non-constraint.

THEORY IN DEPTH · PEIRCE; LIPTON (2004)

Diagnosis is inference to the best explanation

THE CORE IDEA

Abduction: from surprising facts, infer the explanation that would best account for them.

HOW IT WORKS

List rival explanations, ask what each predicts, and test the prediction that separates them.

LIMITS AND CRITIQUE

The best available explanation can still be wrong; it is a working hypothesis.

AT GREENLEAF

High spreadsheet use with the lowest satisfaction is surprising. Which explanation predicts both at once?

THEORY IN DEPTH · KAHNEMAN (2011); KLEIN (2007)

We confirm what we already believe

THE CORE IDEA

Confirmation bias: we search for and weigh evidence that supports our first idea.

HOW IT WORKS

The premortem: imagine the plan has failed a year from now, and write down why. It surfaces hidden risks.

LIMITS AND CRITIQUE

Debiasing techniques help, but cannot remove bias; structure and dissent work better than willpower.

AT GREENLEAF

Before submitting, run a premortem on your diagnosis: if it is wrong, what will the data show in Round 3?

THEORY IN DEPTH · ARGYRIS AND SCHÖN (1978)

Single-loop fixes the action; double-loop fixes the assumption

THE CORE IDEA

Single-loop learning corrects errors within existing goals. Double-loop learning questions the goals themselves.

HOW IT WORKS

Ask not only whether we did it right, but whether we were solving the right problem.

LIMITS AND CRITIQUE

Organizations resist double-loop learning because it threatens established routines and status.

AT GREENLEAF

Buying a new HRIS is single-loop. Asking whether HR should serve the frontline first is double-loop.

Three things to carry out of this room

1

A framework names a pattern. Only you can name a cause.

List three candidates, find the evidence that tells them apart, and say what would prove you wrong. Elimination, not confirmation.

2

Not everything wrong is binding.

The binding constraint is the one whose relaxation changes something downstream you can name. Work away from it and you buy activity, not outcomes.

3

Contradictions are the best evidence you will find.

Explain a disagreement before you reconcile it, and check that your independent sources could not fail together.

READING AND BRIDGE

Where to take this next

Revisit from L1

The six precedence rules, now used as sequencing constraints rather than as descriptions. Unit-1 Core holds Westerman and Kavanagh; Unit-1 Supplementary holds your Week 1 note, the DPDP Act and Fitz-enz & Mattox.

Not in the folder — this session is your source

Goldratt on the theory of constraints · Denzin on triangulation · the clinical literature on premature closure and diagnostic error. Read the slide notes; a short primer will follow.

Practice before the next session

Reopen your Round 1 and Round 2 rationales together. Rewrite the central claim of each in the six-move form — including the candidates you did not consider at the time, and the evidence that would prove you wrong.

Coming next

Unit II — the architecture a diagnosis leads to: the HRIS as system of record, the ATS, the LMS, and the integration layer that decides whether any of them speak.

Lewin and the full eight steps of Kotter come in the change-management session. The fairness critique of predictive models comes in Unit III.